

AI Use Case Intake Checklist

A practical step-by-step checklist for capturing the right information before an AI use case is submitted, triaged, approved, or sent back for cleanup.

Purpose

Use this checklist to make sure the intake process covers the business objective, AI tool profile, data exposure, decision impact, risks, required controls, and approval readiness. It is designed for SMB governance teams that want discipline without red tape.

Document Type	Operational Checklist Template	Version	1.0
Prepared For	{{COMPANY_NAME}}	Owner	{{AI_GOVERNANCE_LEAD}}
Review Cycle	At least annually or upon process change	Related Artifact	AI Use Case Registry
Primary Users	Business owners, AI requestors, IT, security, privacy, legal	Status	Draft / Approved / In Use

How to use this checklist

- Complete the checklist during the intake process - not after the use case is already live.
- If an item cannot be completed, document the gap and identify who must close it before review.
- Use the Notes / Evidence column for links, supporting documents, screenshots, or contract references.
- Any answered risk trigger should be routed into the AI risk triage workflow and the AI Use Case Registry.

1. Initiate the request

Confirm that the use case has a real sponsor, business purpose, and enough context to justify intake.

Check	Required step or confirmation	What should be covered	Notes / evidence
☐	Confirm a named business owner is sponsoring the use case.	Record the owner, department, and who is accountable for outcomes, not just for submitting the request.	
☐	State the business problem or opportunity clearly.	Capture what the team is trying to improve, reduce, automate, or accelerate.	
☐	Define the requested timeline and urgency.	Note expected go-live, pilot, or evaluation dates and why timing matters.	
☐	Confirm whether this is new, existing, or materially changed AI activity.	Identify if the request is a first-time use case, an expansion, or a change to something already in use.	

2. Define the use case

Make the use case understandable enough that reviewers know what the AI is supposed to do and what it should never do.

Check	Required step or confirmation	What should be covered	Notes / evidence
☐	Describe the AI-enabled task or workflow.	Summarize the exact activity, such as drafting, summarizing, classification, chatbot assistance, decision support, or automation.	
☐	Clarify the intended users and affected teams.	Identify who will use the AI, who depends on the output, and which function owns the process.	
☐	Define the expected outputs.	Note whether the output is text, recommendations, analysis, code, customer responses, scoring, or automation triggers.	
☐	Define success measures.	Capture how value will be judged: time saved, cycle-time reduction, quality improvement, revenue lift, error reduction, or user experience improvement.	

3. Profile the tool, model, or vendor

Reviewers should know what technology is involved before they decide whether the request is routine or risky.

Check	Required step or confirmation	What should be covered	Notes / evidence
☐	Identify the AI tool, model, or vendor by name.	List the product, vendor, model family, embedded feature, API, or internal solution involved.	
☐	Confirm the deployment model.	Indicate whether the use case relies on a public consumer tool, enterprise tenant, embedded SaaS feature, direct API, or custom build.	
☐	Check whether the tool is already approved.	Record current approval status and whether an exception is being requested.	
☐	Capture third-party documentation availability.	Confirm whether security papers, privacy terms, DPAs, or product documentation are available if the tool is externally provided.	

4. Understand the data profile

Most intake failures happen here. If the team cannot explain the data, the review should not pretend otherwise.

Check	Required step or confirmation	What should be covered	Notes / evidence
☐	List the data inputs the AI will receive.	Identify prompt content, file uploads, datasets, records, or integrations used by the system.	
☐	Identify the highest data sensitivity involved.	Classify whether the use case touches public, internal, confidential, regulated, customer, employee, financial, or other sensitive information.	
☐	Confirm whether prohibited or restricted data is involved.	Call out any data types that require an approved platform, extra controls, or outright prohibition.	
☐	Check how data is stored, retained, or used for model training.	Capture known vendor terms for retention, reuse, training, logging, and deletion.	

Common failure point

If the requestor cannot explain what data goes in, what outputs come out, and whether the vendor trains on that data, the intake is not complete. That is not a minor detail. That is the review.

5. Assess impact and decision use

The big question is not whether AI is involved. It is whether AI can influence decisions that matter.

Check	Required step or confirmation	What should be covered	Notes / evidence
☐	Determine whether the output informs or automates a decision.	Clarify whether people only reference the output, act on it directly, or allow it to trigger business steps automatically.	
☐	Identify who may be affected by the output.	Note whether customers, employees, applicants, vendors, patients, or members of the public are impacted.	
☐	Confirm whether the use case is internal-facing or external-facing.	Customer- or public-facing use cases typically require tighter review and stronger controls.	
☐	Identify any regulated or high-consequence context.	Flag HR, finance, healthcare, legal, safety, eligibility, credit, or other areas where	

Check	Required step or confirmation	What should be covered	Notes / evidence
		bad output can travel fast and hurt harder.	

6. Run the initial risk screen

This is the fast filter that tells the governance team whether the request is simple, sensitive, or likely to need escalation.

Check	Required step or confirmation	What should be covered	Notes / evidence
☐	Check for sensitive or regulated data exposure.	If yes, route the request for deeper privacy, security, and control review.	
☐	Check for customer-facing or employee-impacting usage.	If yes, document how outputs will be reviewed and what safeguards exist.	
☐	Check for automated action or significant decision support.	If yes, require clearer oversight, validation, and escalation criteria.	
☐	Check for integration with core systems or large-scale deployment.	If yes, confirm architecture, access controls, logging, and rollback readiness.	

Escalation note

Any yes-answer tied to sensitive data, customer or employee impact, automated decisions, or regulated use should trigger deeper review. The checklist is a front door, not a loophole.

7. Confirm required controls and safeguards

A use case should not move forward just because it sounds useful. Useful can still be reckless if guardrails are missing.

Check	Required step or confirmation	What should be covered	Notes / evidence
☐	Define who reviews outputs before action is taken.	Identify the human oversight point, exception thresholds, and escalation path for suspect outputs.	
☐	Confirm access and least-privilege controls.	Document who can use the tool, who can configure it, and who can approve changes.	
☐	Confirm output validation or quality checks.	List any spot checks, sample reviews, testing, or accuracy thresholds used before trust is granted.	
☐	Confirm user awareness or training needs.	State whether employees need policy reminders, do-not-enter guidance, or prompt/output review training before rollout.	

8. Assemble the supporting evidence

A clean intake package saves time, prevents endless follow-up, and makes the registry look like a governance system instead of a junk drawer.

Check	Required step or confirmation	What should be covered	Notes / evidence
☐	Attach or link supporting documentation.	Examples include screenshots, process notes, architecture diagrams, contract documents, DPAs, vendor whitepapers, or pilot results.	
☐	Summarize known risks, concerns, or requested exceptions.	Do not hide sharp edges. Capture them early so the review can address them honestly.	
☐	Document dependencies and affected systems.	List connected applications, data sources, business processes, and support teams.	
☐	Prepare a short reviewer summary.	Provide a quick synopsis of what the use case does, why it matters, and what needs a decision.	

9. Prepare the triage handoff

By this point the request should be reviewable, categorizable, and ready for the next governance step.

Check	Required step or confirmation	What should be covered	Notes / evidence
☐	Propose an initial risk rating.	Estimate whether the request appears low, medium, or high risk based on data, impact, and control complexity.	
☐	Identify which review teams are required.	Mark whether security, privacy, legal, procurement, compliance, or IT operations should be included.	
☐	Confirm the record is ready for the AI Use Case Registry.	The title, owner, tool, status, and next step should be clear enough to log immediately.	
☐	Identify the next workflow state.	Examples: ready for standard review, fast-track review, remediation required, vendor review pending, or escalation needed.	

10. Submission readiness and handoff

Use this final checkpoint before the use case is logged as complete and sent for review, approval, or escalation.

Check	Required step or confirmation	What should be covered	Notes / evidence
☐	A complete intake package has been assembled and obvious gaps have been closed.	All required sections above are complete or the missing items have a named owner and due date.	
☐	The business owner has reviewed the described use case and accepts accountability for its intended use.	There is a named owner for the use case, not just a curious employee with a browser tab and optimism.	
☐	The request has been routed to the right reviewers.	Security, privacy, legal, procurement, or IT reviewers have been engaged where triggers require them.	
☐	The use case is ready to be entered or updated in the AI Use Case Registry.	Registry ID, title, owner, tool/vendor, risk rating, and status are ready for recording.	
☐	A next decision has been defined.	Mark whether the use case is ready for fast-track approval, standard review, deeper assessment, or rejection/remediation.	

Prepared by	{{REQUESTOR_NAME}}	Date	{{DATE}}
Reviewed by	{{REVIEWER_NAME}}	Decision	{{APPROVED / NEEDS MORE INFO / ESCALATE}}
Registry ID	{{REGISTRY_ID}}	Next Review Date	{{DATE}}